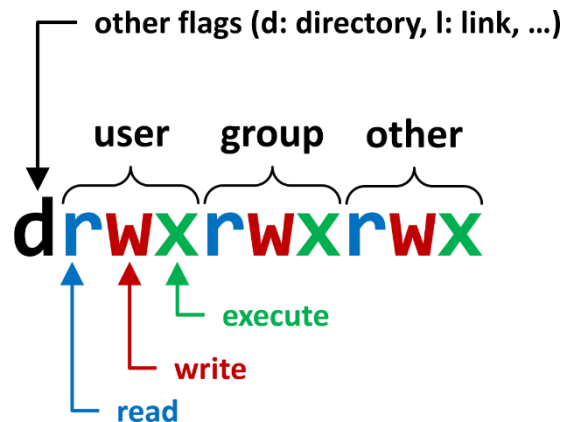



## Input/Output

# Unix Filesystem Concepts



# Module Outline

- The Unix File Concept
- The Unix Filesystem
- Filesystems and Security
- Module Summary



*I/O device abstraction*

# The Unix File Concept

# Unix Files

## ■ A Unix file is a sequence of m bytes:

- $B_0, B_1, B_2, \dots, B_k, \dots, B_{m-1}$

*file system translate this bunch of bytes to meaning for information format type*

$B_0$	$B_1$	$B_2$	...	$B_k$	...	$B_{m-1}$
-------	-------	-------	-----	-------	-----	-----------

- Example: File containing the lower-case English alphabet a-z

0	1	2		12		25
a	b	c	...	m	...	z

# Unix Files

■ in \*nix systems, **everything** is modeled as a file

- file system
- disk, disk partitions
- memory
- USB devices: keyboard, mouse, ...
- display
- network
- audio
- video
- ...

ls

# Unix Files: Examples

file 72 64 471. (mt system)

- I/O devices are represented as files: */dev*

- */dev/sda* (first disk on SATA bus)
- */dev/input/mice* (aggregate of all connected mice)
- */dev/tty2* (terminal)

) *cpu, input...*

- The kernel is exposed with a number of files

- */dev/kmem* (kernel memory image)
- */proc* (*process*) (kernel data structures) : *virtual file system (kernel data struc)*  
*not a real file...*

- System configuration is mapped as files

- */sys*

*cat mem | hexdump -C (less)*

*maps (memory address virtual)*

# Unix File Types

## ■ Regular file

- File containing user/app data (binary, text, whatever)
- OS does not know anything about the format
  - ▶ other than “sequence of bytes”, akin to main memory

## ■ Directory file

- A file that contains the names and locations of other files

## ■ Character special and block special files

- Terminals (character special) and disks (block special)

↳ unit of byte (char)

unit of block not bytes (char)

## ■ FIFO (named pipe)

- A file type used for inter-process communication

## ■ Socket

- A file type used for (local or networked) communication between processes

# Unix I/O

## ■ Key Features

- Design concept: All input and output is handled in a consistent and uniform way
- Elegant mapping of files to devices allows kernel to export simple interface called Unix I/O

**One single file interface to interact with any kind of device**

(well, almost)



# The Unix Filesystem

# The Unix Filesystem

## ■ One root to rule them all

- single file system starting with at the root ("/")
- unlike Windows, there is no concept of a “drive”
- additional filesystems are mapped into the file system tree as a directory
- mount point = directory where a filesystem is attached

```
devel@csapvm $ ls -l
share
temp
work
devel@csapvm $ mkdir extern
devel@csapvm $ echo "hello" > extern/hello
devel@csapvm $ ls extern/
hello
devel@csapvm $ sudo mount -t tmpfs /dev/shm extern/
Password:
devel@csapvm $ ls extern/
extern
devel@csapvm $ echo "extern" > extern/extern
devel@csapvm $ ls extern/
extern
devel@csapvm $ sudo umount extern
devel@csapvm $ ls extern/
hello
devel@csapvm
```

# The Unix Filesystem

What is mount

## ■ One root to rule them all

- mapped filesystems hide the contents of the directory tree under the mount point

```
devel@csapvm $ ls -l
share
temp
work
devel@csapvm $ mkdir extern
devel@csapvm $ echo "hello" > extern/hello
devel@csapvm $ ls extern/
hello
devel@csapvm $ sudo mount -t tmpfs /dev/shm extern/
Password:
devel@csapvm $ ls extern/
extern
devel@csapvm $ echo "extern" > extern/extern
devel@csapvm $ ls extern/
extern
devel@csapvm $ sudo umount extern
devel@csapvm $ ls extern/
hello
devel@csapvm
```

→ repeat

→ create dir

Link holder with others...

device: /dev/sda4

```
/
|-home
|-devel
|-share
|-temp
|-work
|-extern
|-hello
```

device: host (vboxsf)

```
/
|-grades
|-labs
|-lectures
|-resources
|-videos
```

device: /dev/shm (tmpfs)

```
/
|-home
|-devel
|-share
|-temp
|-work
|-extern
|-hello
```

```
/
|- extern
```

# The Unix Filesystem

## ■ One root to rule them all

- extremely powerful concept
- each mounted filesystem can have additional properties *(Auth)*
  - ▶ do not allow writes (ro)
  - ▶ do not update access time (noatime)
  - ▶ do not allow execution of programs (noexec)
  - ▶ do not allow set user/group id (nosuid)
  - ▶ ...

```
devel@csapvm $ mount
/dev/sda4 on / type ext4 (rw,noatime)
...
none on /tmp type tmpfs (rw,noatime,size=262144k)
none on /var/tmp type tmpfs (rw,noatime,size=131072k)
devel_share on /home/devel/share type vboxsf (rw,nodev,relatime,ioccharset=utf8,uid=1000,gid=100)
...
```

# The Unix Filesystem

## ■ User, group, soft & hard links

- each file is owned by a *user* and a *group*

bernhard@intel72:/home\$ ls *Long format*

```
total 148
drwxr-xr-x 13 root      root    4096 Feb  4 2017 alumni
drwxr-xr-x  4 anna      student 4096 Aug 31 2016 anna
-rw----- 1 root      root    8192 Jan 20 2017 aquota.user
drwxr-xr-x  8 barend    student 4096 Dec 19 2018 barend
drwxr-xr-x 16 bernhard  prof   4096 Sep 16 00:47 bernhard
drwxr-xr-x  7 camilo    student 4096 Jun 28 2017 camilo
drwxr-xr-x  6 changmin  student 4096 Mar 23 2017 changmin
drwxr-xr-x 13 changyeon postdoc 4096 Feb 19 2021 changyeon
...
drwxr-xr-x  4 simeon    student 4096 Mar 10 2017 simeon
drwxr-xr-x  5 root      root    4096 Apr 21 2020 svn
drwxr-xr-x  6 vm        users   4096 May 10 10:43 vm
drwxr-xr-x  5 vmmail    vmmail  4096 Apr 19 2012 vmmail
drwxr-xr-x  4 xinyi     student 4096 Aug 31 2016 xinyi
drwxr-xr-x 20 younghyun postdoc 4096 Jun 18 13:56 younghyun
drwxr-xr-x  8 youngsu   student 4096 Jun 17 2019 youngsu
```

type/  
permissions

user / group

hard links

size (bytes)

modification time

file name

# The Unix Filesystem

*→ pointer/url...*

## ■ User, group, soft & hard links

- files can point to each other via soft or hard links

```
$ printf "hello" > hello.txt
$ printf "test" > test.txt
$ ls -l
total 8
-rw-r--r-- 1 devel users 5 Sep 27 01:10 hello.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:11 test.txt

$ ln hello.txt helloworld.txt
$ ln -s test.txt testing
$ ls -l
total 12
-rw-r--r-- 2 devel users 5 Sep 27 01:10 hello.txt
-rw-r--r-- 2 devel users 5 Sep 27 01:10 helloworld.txt
lrwxrwxrwx 1 devel users 8 Sep 27 01:12 testing -> test.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:11 test.txt
```

*count of hard link*


**hard link count**

**1: soft (symbolic) link flag**

**In: create hard link**

**In -s: create soft (symbolic) link**

*kind of 다른거?*



# The Unix Filesystem

## ■ User, group, soft & hard links

- files can point to each other via soft or hard links

```
$ printf "hello" > hello.txt
$ printf "test" > test.txt
$ ls -l
total 8
-rw-r--r-- 1 devel users 5 Sep 27 01:10 hello.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:11 test.txt
```

```
$ ln hello.txt helloworld.txt
$ ln -s test.txt testing
```

```
$ ls -l
total 12
-rw-r--r-- 2 devel users 5 Sep 27 01:12 hello.txt
-rw-r--r-- 2 devel users 5 Sep 27 01:12 helloworld.txt
lrwxrwxrwx 1 devel users 8 Sep 27 01:12 testing -> test.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:11 test.txt
```

...

```
$ printf ", world!"\\n >> helloworld.txt
```

```
$ cat helloworld.txt
```

```
hello, world!
```

```
$ cat hello.txt
```

```
hello, world!
```

```
$ ls -l
```

```
total 12
```

```
-rw-r--r-- 2 devel users 14 Sep 27 01:12 hello.txt
```

```
-rw-r--r-- 2 devel users 14 Sep 27 01:12 helloworld.txt
```

```
lrwxrwxrwx 1 devel users 8 Sep 27 01:12 testing -> test.txt
```

```
-rw-r--r-- 1 devel users 4 Sep 27 01:11 test.txt
```

*Analysis ls -la*

*Any At home.*

*copy.*

# The Unix Filesystem

## ■ User, group, soft & hard links

- files can point to each other via soft or hard links

```
$ printf "hello" > hello.txt
$ printf "test" > test.txt
$ ls -l
total 8
-rw-r--r-- 1 devel users 5 Sep 27 01:10 hello.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:11 test.txt
```

```
$ ln hello.txt helloworld.txt
$ ln -s test.txt testing
```

```
$ ls -l
total 12
-rw-r--r-- 2 devel users 5 Sep 27 01:12 hello.txt
-rw-r--r-- 2 devel users 5 Sep 27 01:12 helloworld.txt
lrwxrwxrwx 1 devel users 8 Sep 27 01:12 testing -> test.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:13 test.txt
```

...

```
$ printf ", test, and test!"\n >> testing
```

```
$ cat test.txt
```

```
test, test, and test!
```

```
$ cat testing
```

```
test, test, and test!
```

```
$ ls -l
```

```
total 12
```

```
-rw-r--r-- 2 devel users 14 Sep 27 01:12 hello.txt
```

```
-rw-r--r-- 2 devel users 14 Sep 27 01:12 helloworld.txt
```

```
lrwxrwxrwx 1 devel users 8 Sep 27 01:12 testing -> test.txt
```

```
-rw-r--r-- 1 devel users 22 Sep 27 01:13 test.txt
```



# The Unix Filesystem

## ■ User, group, soft & hard links

- files can point to each other via soft or hard links

```
$ printf "hello" > hello.txt
$ printf "test" > test.txt
$ ls -l
total 8
-rw-r--r-- 1 devel users 5 Sep 27 01:10 hello.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:11 test.txt
```

```
$ ln hello.txt helloworld.txt
$ ln -s test.txt testing
```

```
$ ls -l
total 12
-rw-r--r-- 2 devel users 5 Sep 27 01:12 helloworld.txt
-rw-r--r-- 2 devel users 5 Sep 27 01:12 test.txt
lrwxrwxrwx 1 devel users 8 Sep 27 01:12 testing -> test.txt
-rw-r--r-- 1 devel users 4 Sep 27 01:12 test.txt
```

...

```
$ rm hello.txt
$ rm test.txt
```

```
$ ls -l
total 4
-rw-r--r-- 1 devel users 14 Sep 27 01:12 helloworld.txt
lrwxrwxrwx 1 devel users 8 Sep 27 01:12 testing -> test.txt
```

```
$ cat helloworld.txt
hello, world!
```

```
$ cat testing
cat: testing: No such file or directory
```

**broken (soft) link**

# The Unix Filesystem

## ■ File types

- Unix knows different types of files

```
/tmp/demo $ ls -l
total 8
drwxr-xr-x  2 bernhard  users  40 Sep  3 20:00 directory
prw-r--r--  1 bernhard  users   0 Sep  3 20:01 fifo
-rw-r--r--  2 bernhard  users  14 Sep  3 20:03 file
-rw-r--r--  2 bernhard  users  14 Sep  3 20:03 hardlink
lrwxrwxrwx  1 bernhard  users   4 Sep  3 20:00 softlink -> file
```

↑  
file type

Letter	File type
-	regular file
d	directory
l	soft link (symbolic link)
p	named pipe (fifo)
s	socket
c	character device file (see /dev)
b	block device file (see /dev) <i>only I/O with unit of block.</i>

# The Unix Filesystem

## ■ Hidden files

- Hidden files start with a “.” in Unix file systems

```
/tmp/demo $ ls -la
total 8
drwxr-xr-x  4 bernhard users 180 Sep  3 20:03 .
drwxrwxrwt 18 root     root 1120 Sep  3 19:59 ..
drwxr-xr-x  2 bernhard users  40 Sep  3 20:00 directory
prw-r--r--  1 bernhard users   0 Sep  3 20:01 fifo
-rw-r--r--  2 bernhard users  14 Sep  3 20:03 file
-rw-r--r--  2 bernhard users  14 Sep  3 20:03 hardlink
drwxr-xr-x  2 bernhard users  40 Sep  3 20:00 .hiddendir
-rw-r--r--  1 bernhard users   0 Sep  3 20:00 .hiddenfile
lrwxrwxrwx  1 bernhard users   4 Sep  3 20:00 softlink -> file
```

- Special (hidden) entries
  - ▶ “.” current directory
  - ▶ “..” parent directory (hence “cd ..”)

# Filesystem Hierarchy Standard (FHS)

## ■ What goes where?

- Is it all a big mess?

```
$ $ ls -l /
total 76
drwxr-xr-x  2 root root  4096 Sep  3 02:11 bin
drwxr-xr-x  3 root root  4096 Aug  1 20:37 boot
drwxr-xr-x 20 root root 4020 Sep  1 23:17 dev
drwxr-xr-x 88 root root  4096 Sep  3 11:59 etc
drwxr-xr-x  5 root root  4096 Mar 24  2022 home
drwxr-xr-x 13 root root  4096 Jul  5 21:40 lib
drwxr-xr-x  7 root root  4096 Sep  3 02:11 lib64
drwx----- 2 root root 16384 Mar 23  2022 lost+found
drwxr-xr-x  2 root root  4096 Mar 21  2022 media
drwxr-xr-x  4 root root  4096 Sep 16  2022 mnt
drwxr-xr-x 16 root root  4096 Jul 16 19:05 opt
dr-xr-xr-x 447 root root      0 Aug 19 21:52 proc
drwx----- 7 root root  4096 Sep  3 11:59 root
drwxr-xr-x 19 root root   740 Sep  1 23:16 run
drwxr-xr-x  2 root root 12288 Sep  3 02:11 sbin
dr-xr-xr-x 12 root root      0 Aug 19 21:52 sys
drwxrwxrwt 17 root root  1100 Sep  3 18:51 tmp
drwxr-xr-x 12 root root  4096 Sep 20  2022 usr
drwxr-xr-x 10 root root  4096 Sep  3 01:37 var
```

# Filesystem Hierarchy Standard (FHS)

## ■ The Filesystem Hierarchy Standard (FHS)

- Conventions for the layout of directories and files on Unix systems
- Maintained by the Linux Foundation: <https://refspecs.linuxfoundation.org/fhs.shtml>
  - ▶ initial version released February 14, 1994, current version: 3.0 (June 3, 2015)

Directory	Description
/	<b>root</b>
/bin	Essential binaries required during boot-up
/boot	boot loader, kernel
/dev	Device files (disks, partitions, memory, audio, video, ...)
/etc	Host-specific, system-wide configuration files
/home	User home directories
/lib[64]	System libraries (required by binaries in /bin, /sbin)
/media, /mnt	Mount points for removable media
...	

# Filesystem Hierarchy Standard (FHS)

## ■ The Filesystem Hierarchy Standard (FHS)

- <https://refspecs.linuxfoundation.org/fhs.shtml>

Directory	Description
...	
/opt	Additional application software packages
/proc	Process and kernel information (virtual filesystem)
/root	Home directory for the root user (administrator)
/run	Run-time variable data
/sbin	Essential system binaries
/sys	Device driver & kernel information and configuration
/tmp, /var/tmp	Temporary files (often not preserved across reboots)
/usr	Secondary hierarchy for read-only user data
/var	Variable files

# Filesystem Hierarchy Standard (FHS)

## ■ The Filesystem Hierarchy Standard (FHS)

- <https://refspecs.linuxfoundation.org/fhs.shtml>

Directory	Description
<b>/usr</b>	<b>Secondary hierarchy for read-only user data</b>
/usr/bin, /usr/sbin	Non-essential binaries
/usr/include	Standard include files (C headers)
/usr/lib[64]	Libraries required by binaries in /usr/bin, /usr/sbin
/usr/libexec	Binaries run via scripts (do not run directly)
/usr/local	Tertiary hierarchy for local data specific to this machine
/usr/share	Applicate shared data
/usr/src	Source code (kernel source)

# Filesystem Hierarchy Standard (FHS)

## ■ The Filesystem Hierarchy Standard (FHS)

- <https://refspecs.linuxfoundation.org/fhs.shtml>

Directory	Description
<b>/var</b>	<b>Variable files</b>
/var/cache	Application cache data
/var/db	Gentoo portage (config & source files)
/var/lib	Persistent state data modified by applications
/var/lock	Lock files to keep track of resources currently in use
/var/log	System log files
/var/mail	Mailbox files (servers only)
/var/run	Run-time variable data (FHS 3.0: mapped to /run)
/var/spool	Spool for tasks waiting to be processed
/var/tmp	Temporary files (often not preserved across reboots)



# Filesystem Hierarchy Standard (FHS)

## ■ The Filesystem Hierarchy Standard (FHS)

- <https://refspecs.linuxfoundation.org/fhs.shtml>

Directory	Description
<u>/home/&lt;USER&gt;/</u>	★ <u>User-specific files</u>
.cache	User cached data
.config	User-specific configuration
.local	User local data (bin, lib, share)
.mozilla	Application-specific configuration, z.B. for Mozilla
.ssh	idem for SSH
.vim	and for vim
	plus user-created directories and files

- Note: “~” is an abbreviation for /home/<USER>
  - ▶ ls ~/.config

```
fstatat(dd, entry->d_name, &sb, AT_SYMLINK_NOFOLLOW);

if (S_ISREG(sb.st_mode) &&
    (((sb.st_uid == 0) && (sb.st_mode & S_ISUID)) ||
     ((sb.st_gid == 0) && (sb.st_mode & S_ISGID)))
{
    fstatfs(dd, &dsb);
    if (!(dsb.f_flags & (ST_NOEXEC|ST_NOSUID))) {
        // dangerous configuration
    }
}
```

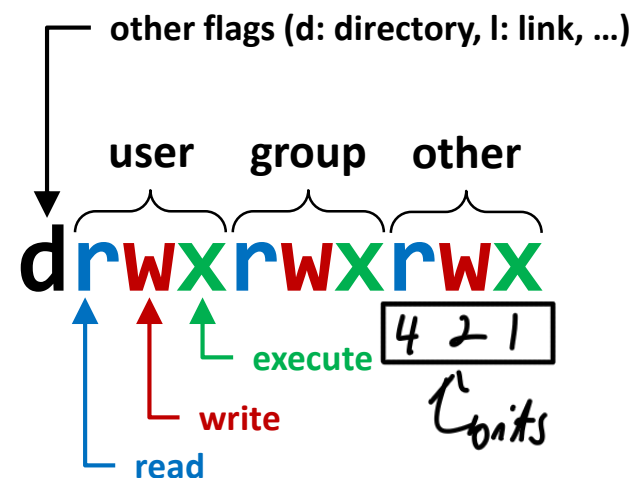
# Filesystems and Security

# Filesystems and Security

## ■ Standard \*nix Access Control Lists (ACL)

- three levels of access
  - ▶ owner, group, other
- three kinds of permissions
  - ▶ read (r), write (w), execute (x)

```
devel@csapvm $ ls -l
total 60
drwxr-xr-x 2 devel devel 4096 Sep 26 19:44 doc
-rwxr-xr-- 1 devel devel 18440 Sep 27 01:50 lsec
-rw-r--r-- 1 devel devel 9842 Sep 27 00:52 lsec.c
-rw-r--r-- 1 devel devel 5493 Sep 27 00:37 lsec.h
-rw-rw-rw- 1 devel devel 1009 Sep 26 19:44 Makefile
-r--r----- 1 devel devel 22738 Sep 26 19:44 README.md
drwxr-xr-x 2 devel devel 4096 Sep 26 19:44 reference
drwx---r-x 2 devel devel 4096 Sep 26 19:44 tools
```



$r - - = 4$   
 $r - x = 5 (4+1)$

- eXecute permission on
- file: execute (run) program
  - directory: list contents of directory

# Filesystems and Security

## ■ Standard \*nix Access Control Lists

- modify with chmod command

```
devel@csapvm ~/temp $ echo "hello" > test.txt
devel@csapvm ~/temp $ ls -l
-rw-r--r-- 1 devel devel 6 Sep 16 19:37 test.txt

devel@csapvm ~/temp $ chmod g+w test.txt
devel@csapvm ~/temp $ ls -l
-rw-rw-r-- 1 devel devel 6 Sep 16 19:37 test.txt

devel@csapvm ~/temp $ echo "ls -l" > script.sh
devel@csapvm ~/temp $ ./script.sh
-bash: ./script.sh: Permission denied
devel@csapvm ~/temp $ ls -l
-rw-r--r-- 1 devel devel 6 Sep 16 19:38 script.sh

devel@csapvm ~/temp $ chmod 750 script.sh
devel@csapvm ~/temp $ ls -l
-rwxr-x--- 1 devel devel 6 Sep 16 19:37 test.txt
```

group other (u, g, w)

man  
manual

- for details, see man 1 chmod

# Filesystems and Security

## ■ Security-related settings

- Sticky bit *(not super important)*

```
devel@csapvm $ ls -ld /tmp  
drwxrwxrwt 6 root root 240 Oct  8 14:32 /tmp
```

- ▶ files in a directory with the sticky bit set can be renamed or deleted only by the owner of the file, by the owner of the directory, and by a privileged user/process
  - should be set on world-writable directories such as /tmp
- ▶ flag available in struct stat (S\_ISVTX)
  - man 2 stat
  - man 7 inode

# Filesystems and Security

## ■ Security-related settings

- Set owner User ID upon execution (SUID/SGID bit)

```
develop@csapvm $ ls -l /usr/bin/sudo  
-rws--x--x 1 root root 204320 Sep  1 16:17 /usr/bin/sudo
```

- ▶ binary executed with permissions of owner (root in this case) (as opposed to context of user who executes it)
  - useful to give temporary permissions of owner
  - only trusted binaries owned by root must have `suid/sgid` bits set!
- ▶ flag available in struct `stat` (`S_ISUID` / `S_ISGID`): `man 2 stat`, `man 7 inode`
- ▶ filesystem may disallow SUID:  
use `fstatfs()` and check for `ST_NOSUID` flag

# Filesystems and Security

*skipped*

## ■ Security-related settings

- World-writable directory with execute permission

```
devel@csapvm $ ls -ld /tmp
drwxrwxrwx 6 root root 240 Oct  8 14:32 /tmp
or
drwxrwxrwt 6 root root 240 Oct  8 14:32 /tmp
```

- ▶ world-writable directories on a file system with execute permission are a security risk
- ▶ anyone with access to the system may place an executable and run it (typical scenario: webserver breach → write script to /tmp → execute it)
- ▶ world-writable flag in struct stat (S\_IWOTH)
- ▶ filesystem may disallow execution:  
use `fstatfs()` and check for `ST_NOEXEC` flag

# Example: Explore SUID/GUID

- Set user/group owner  
`sudo chown <usr>[:<grp>]`

- Set suid/sgid bit

## suid bit

`sudo chmod 4755 <exe>`

## sgid bit

`sudo chmod 2755 <exe>`

## suid+sgid bits

`sudo chmod 6755 <exe>`

```
devel@csapvm ~/work/03 $ gcc -o wai whoami.c
devel@csapvm ~/work/03 $ ./wai
User & group information
-----
User:          devel          (1000)
Group:         devel          (1000)

Effective user:  devel          (1000)
Effective group: devel          (1000)

devel@csapvm ~/work/03 $ sudo chown tester:users wai
devel@csapvm ~/work/03 $ sudo chmod 4755 wai
devel@csapvm ~/work/03 $ ls -l
total 20
-rwsr-xr-x 1 tester users 16000 Sep 27 02:32 wai
-rw-r--r-- 1 devel  devel  1380 Sep 27 02:31 whoami.c

devel@csapvm ~/work/03 $ ./wai
User & group information
-----
User:          devel          (1000)
Group:         devel          (1000)

Effective user:  tester        (1001)
Effective group: devel          (1000)
```



# Example: Explore SUID/GUID

## ■ whoami.c: print user/group

- **effective user/group**  
user/group under whose permission the process is executed
- **(real) user/group**  
original user/group under which the process was started (before suid/sgid)

```
#include <grp.h>
#include <pwd.h>

...
int main(int argc, char *argv[]) {
    // get user id, effective user id, group id, effective group id
    uid_t uid = getuid();
    uid_t euid = geteuid();
    gid_t gid = getgid();
    gid_t egid = getegid();

    // get user and group names
    char *user, *euser, *group, *egroup;
    struct passwd *pwd;
    struct group *grp;

    if ((pwd = getpwuid(uid)) != NULL) user = strdup(pwd->pw_name);
    if ((pwd = getpwuid(euid)) != NULL) euser = strdup(pwd->pw_name);
    if ((grp = getgrgid(gid)) != NULL) group = strdup(grp->gr_name);
    if ((grp = getgrgid(egid)) != NULL) egroup = strdup(grp->gr_name);

    // print results
    printf("User & group information\n"
        "-----\n"
        "  User:                %-16s (%4d)\n"
        "  Group:               %-16s (%4d)\n"
        "\n"
        "  Effective user:      %-16s (%4d)\n"
        "  Effective group:     %-16s (%4d)\n",
        user ? user : "n/a", uid,
        group ? group : "n/a", gid,
        euser ? euser : "n/a", euid,
        egroup ? egroup : "n/a", egid);

    // free allocated memory
    free(user); free(egroup); free(group); free(euser);

    return EXIT_SUCCESS;
}
```

whoami.c

# Example: Checking for SUID/SGID Bit

```
DIR *d = opendir(name);
int dd = dirfd(d);
struct dirent *entry;

while ((entry = getNext(d)) != NULL) {
    struct stat sb;
    struct statfs dsb;

    fstatat(dd, entry->d_name, &sb,                // get metadata of directory entry
            AT_SYMLINK_NOFOLLOW);

    if (S_ISREG(sb.st_mode) &&                      // if it's a regular file and
        (((sb.st_uid == 0) && (sb.st_mode & S_ISUID)) || // the user is root & SUID is set or
         ((sb.st_gid == 0) && (sb.st_mode & S_ISGID))) // the group is root & SGID is set
    {
        fstatfs(dd, &dsb);                          // get metadata of file system
        if (!(dsb.f_flags & (ST_NOEXEC|ST_NOSUID))) { // if neither NOEXEC nor NOSUID are set
            // dangerous configuration                // then this is potentially dangerous
        }
    }
}

...
}
```

- for readability, no error checking performed.

# Extended File System Security Concepts



- **POSIX ACLs are limited to access permissions for the user, a group, and everybody else**
- **Extended file attributes (xattrs) provide an extensible and more flexible way to store meta data (including ACLs) about a file**
  - xattrs are key=value pairs where
    - ▶ key has the form “namespace.attribute”
    - ▶ and value is a string
  - currently, xattrs defines the namespaces “security”, “system”, “trusted”, and “user”
    - ▶ security: used by kernel modules such as SELinux to implement advanced ACLs
    - ▶ system: used by the kernel to store system objects
    - ▶ trusted: attributed only visible to processes with the CAP\_SYS\_ADMIN capability
    - ▶ user: store arbitrary additional information about a file such as its mime type, md5sum, character encoding, etc.

# Extended File System Security Concepts

- Check whether a filesystem supports xattrs

```
devel@csapvm $ mount | grep "/"  
/dev/sda4 on / type ext4 (rw,noatime)  
devel@csapvm $ cat /proc/fs/ext4/sda4/options | grep xattr  
user_xattr
```

- Set/get ACLs with setfacl / getfacl

```
devel@csapvm $ echo "Hello" > file.txt  
devel@csapvm $ ls -l file.txt  
-rw-r--r-- 1 devel devel 6 Mar  7 19:75 file.txt  
devel@csapvm $ setfacl -m u:svn:r file.txt  
devel@csapvm $ ls -l file.txt  
-rw-r--r--+ 1 devel devel 6 Mar  7 19:75 file.txt  
devel@csapvm $ getfacl file.txt  
# file: file.txt  
# owner: devel  
# group: devel  
user::rw-  
user:svn:r--  
group::r--  
mask::r--  
other::r--
```

# Extended File System Security Concepts

- Set/get arbitrary attributes with setattr/getattr

```
devel@csapvm $ md5sum file.txt
09f7e02f1290be211da707a266f153b3  file.txt
devel@csapvm $ setattr -n user.checksum.md5 -v
09f7e02f1290be211da707a266f153b3 file.txt
devel@csapvm $ ls -l
total 8
-rw-r--r--+ 1 devel devel 6 Mar  7 19:75 file.txt
devel@csapvm $ getfattr file.txt
# file: file.txt
user.checksum.md5

devel@csapvm $ getfattr -n user.checksum.md5 file.txt
# file: file.txt
user.checksum.md5="09f7e02f1290be211da707a266f153b3"
devel@csapvm $ setattr -x user.checksum.md5 file.txt
devel@csapvm $ getfattr file.txt
devel@csapvm $ getfattr -n user.checksum.md5 file.txt
file.txt: user.checksum.md5: No such attribute
```

## Summary

- -----
- -----
- -----
- -----
- -----
- -----

# Module Summary

# Summary

- Unix concept: everything is a file
- Filesystems support many “advanced” features
  - mount points of different filesystems under common root
  - mounting with different permissions
  - hard and soft links
  - set user/group id
- Security
  - Access Control Lists (ACL)
    - ▶ for user, group, and other (=everybody else)
    - ▶ read, write, and execute permission
  - many “dangerous” configurations possible, especially sticky, suid/sgid bits
  - Extended file attributes (xattrs) provide finer-grained control