


Input/Output

Files and Directories



Module Outline

- File Metadata
- Directories
- Kernel File Management
- Module Summary

```
struct stat {  
    dev_t      st_dev;  
    ino_t      st_ino;  
    mode_t     st_mode;  
    nlink_t    st_nlink;  
    uid_t      st_uid;  
    gid_t      st_gid;  
    dev_t      st_rdev;  
    off_t      st_size;  
    blksize_t  st_blksize;  
    blkcnt_t   st_blocks;  
    struct timespec st_atim;  
    struct timespec st_mtim;  
    struct timespec st_ctim;  
};
```

File Metadata

File Metadata

- Metadata is data about data, in this case information about a file
 - filename
 - file type
 - file size
 - file creation, modification, and access time
 - file access permissions
 - ...

- Per-file metadata maintained by kernel
 - not all metadata is stored in the same place
 - ▶ filename: stored in the directory that contains the files
 - ▶ everything else is stored in the inode of the file
 - inode: internal storage unit used by Unix file systems

File Metadata

- File metadata can be accessed with the `*stat*` family of Unix I/O API calls
 - return the file metadata in a struct `stat`

```
struct stat {
    dev_t          st_dev;          // device
    ino_t          st_ino;         // inode
    mode_t         st_mode;        // protection and file type
    nlink_t        st_nlink;       // number of hard links
    uid_t          st_uid;         // user ID of owner
    gid_t          st_gid;         // group ID of owner
    dev_t          st_rdev;        // device id (if special file)
    off_t          st_size;        // total size, in bytes
    unsigned long  st_blksize;     // preferred blocksize for filesystem I/O
    unsigned long  st_blocks;     // number of 512b blocks allocated
    // high-precision timestamps (precision: nanoseconds; since Linux kernel >=2.6)
    struct timespec st_atim;       // time of last access
    struct timespec st_mtim;       // time of last modification
    struct timespec st_ctim;       // time of last status change
    // low-precision timestamps (precision: seconds; always available)
    time_t         st_atime;       // time of last access
    time_t         st_mtime;       // time of last modification
    time_t         st_ctime;       // time of last status change
};
```

File Metadata

■ struct stat

- manual: stat(2), inode(7)
- st_mode encodes the type and access permissions of a file
 - ▶ use S_IS***(stat.st_mode) macros for convenience
if (S_ISREG(stat.st_mode)) // regular file
 - ▶ use S_IR/W/X*** masks to check the access permissions
if (stat.st_mode & S_IRUSR) // owner has read permission
- st_uid / st_gid
 - ▶ use getpwuid(stat.st_uid) / getgrgid(sb.st_gid) to retrieve user/group information
- st_size vs st_blocks
 - ▶ file is sparse if $\text{st_size} / 512 > \text{st_blocks}$

File Metadata

■ struct stat

- timestamps

- ▶ `st_?tim`: nanosecond precision since Linux 2.6
- ▶ `st_?time`: Unix / Linux <2.6 low-resolution timestamps (precision: seconds)
- ▶ traditionally, Unix filesystems did not record a file's creation date (birth date)
 - newer file systems may support it, retrieve via `statx()` system call

- `st_ctim[e]` vs `st_mtim[e]`

- ▶ `ctim[e]`: timestamp of *last update to file inode* (file meta data)
- ▶ `mtim[e]`: timestamp of *last update to file data*
- ▶ in general, $\text{mtim}[e] \leq \text{ctim}[e]$

- `st_atim[e]`

- ▶ time stamp of *last access to file*
- ▶ often disabled for performance reasons (mount filesystem with `-o noatime`)

File Metadata

- C standard library functions to access file metadata (header sys/stat.h)

Operation	API	Remarks
Retrieve file metadata	<code>int stat(const char *pathname, struct stat *statbuf)</code>	
	<code>int lstat(const char *pathname, struct stat *statbuf)</code>	does not follow symbolic links
	<code>int fstat(int fd, struct stat *statbuf)</code>	
	<code>int fstatat(int dirfd, const char *pathname, struct stat *statbuf, int flags)</code>	useful when stat-ing files in a directory
	<code>int statx(int dirfd, const char *pathname, int flags, unsigned int mask, struct statx *statxbuf)</code>	get extended file status


```
| -a
| ` -b
|   | -c
|   | ` -d
|   ` -f
| -dir1
| ` -seventyfive
| -dir2
| | -three
| ` -seven
` -dir3
  | -five
  ` -twentytwo
```

Directories

Directories

- C standard library functions to deal with directories (headers dirent.h, sys/types.h)

Operation	API	Variants
Open	<code>DIR* opendir(const char *name)</code>	<code>fdopendir</code>
Read entry	<code>struct dirent* readdir(DIR *dirp)</code>	
Close	<code>int closedir(DIR *dirp)</code>	
Retrieve descriptor	<code>int dirfd(DIR *dirp)</code>	
Make directory	<code>int mkdir(const char *pathname, mode_t mode)</code>	<code>mkdirat</code>

Example: Accessing Directories and File Metadata

```
#include <...>

int main(int argc, char *argv[])
{
    DIR *dir = opendir(argc > 1 ? argv[1] : ".");
    if (dir == NULL) { perror("Cannot open directory"); return EXIT_FAILURE; }

    int dd = dirfd(dir);


    struct dirent *e;
    errno = 0;
    while ((e = readdir(dir)) != NULL) {
        struct stat sb;
        if (fstatat(dd, e->d_name, &sb, 0) < 0) {
            perror("Cannot stat file");
        } else {
            printf("  %-32s %10ld\n", e->d_name, sb.st_size);
        }
        errno = 0;
    }
    if (errno != 0) perror("Cannot enumerate directory");

    closedir(dir);

    return EXIT_SUCCESS;
}
```

```
$ ./statter /
opt                4096
sys                0
proc               0
devnull            64
usr                4096
lib64              12288
home               4096
.                  4096
data               4096
var                4096
lost+found         16384
..                 4096
dev                4300
...
```

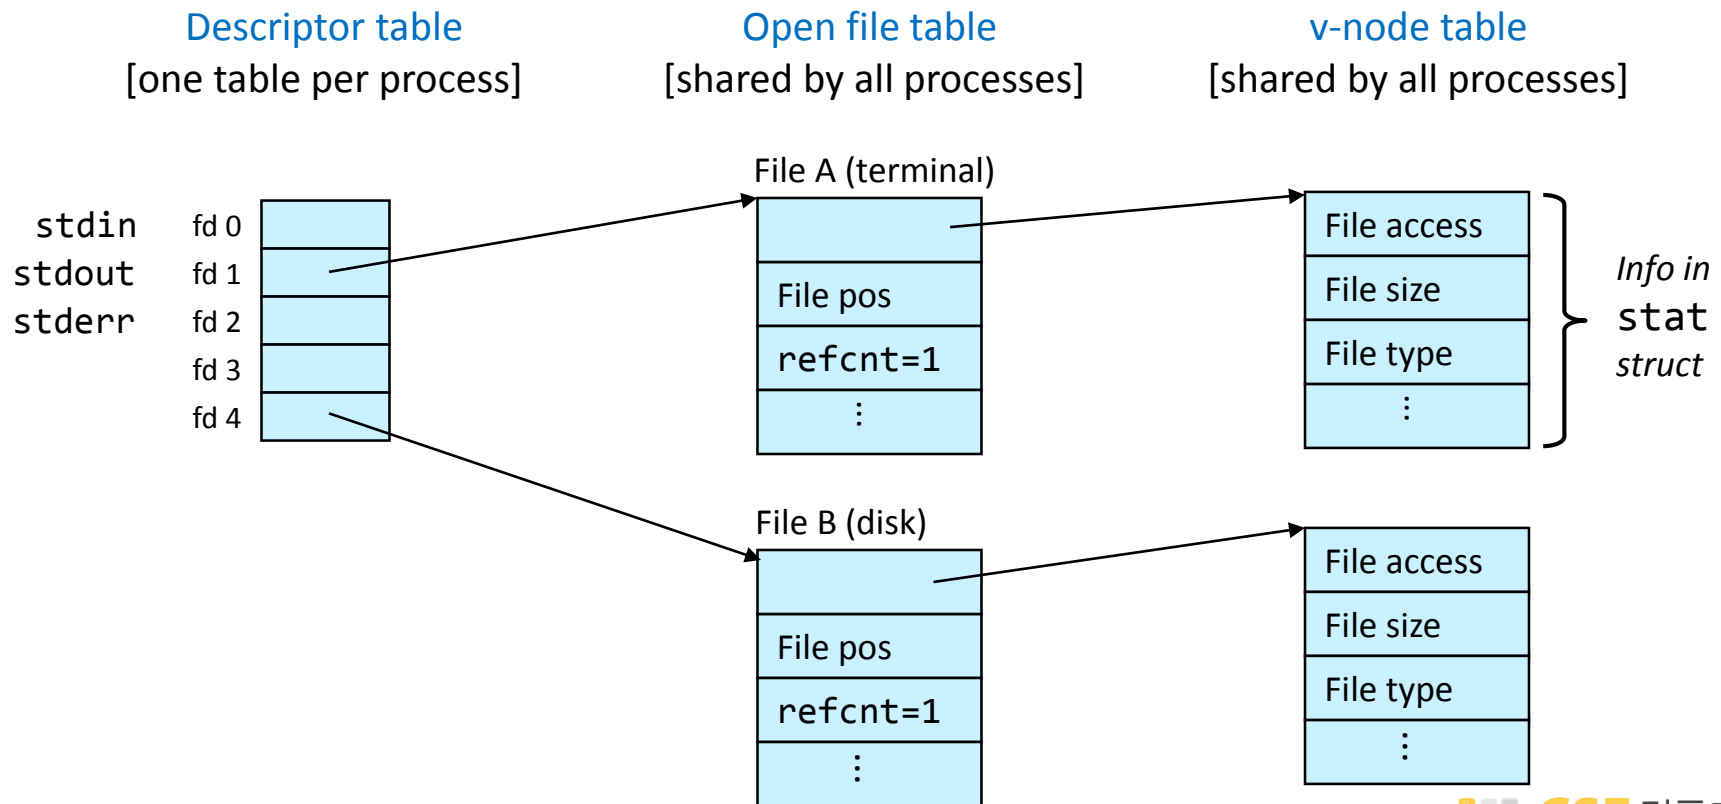
statter.c



Kernel File Management

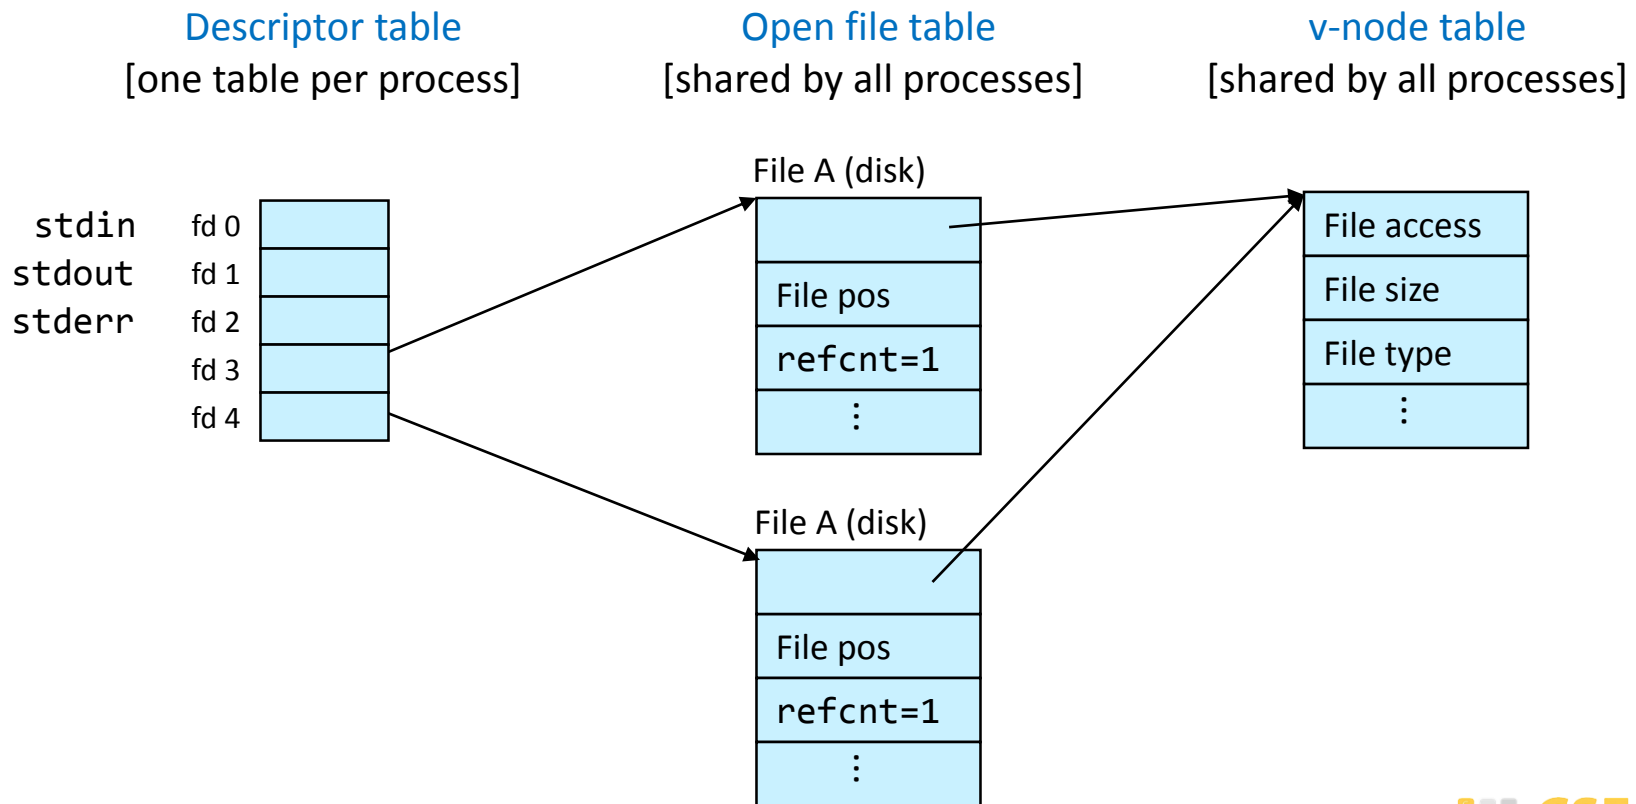
How the Unix Kernel Represents Open Files

- Two descriptors referencing two distinct open disk files. Descriptor 1 (stdout) points to terminal, and descriptor 4 points to open disk file



File Sharing

- Two distinct descriptors sharing the same disk file through two distinct open file table entries
 - E.g., Calling open twice with the same filename argument



I/O Redirection

- I/O redirection is one of the core concepts of Unix

```
$ ls > output.txt
```

```
$ ls | sort -R
```

```
$ cat < input.txt
```

- How does a shell implement I/O redirection?


File Descriptor Manipulation

- C standard library functions allow manipulation of file descriptors (header unistd.h)

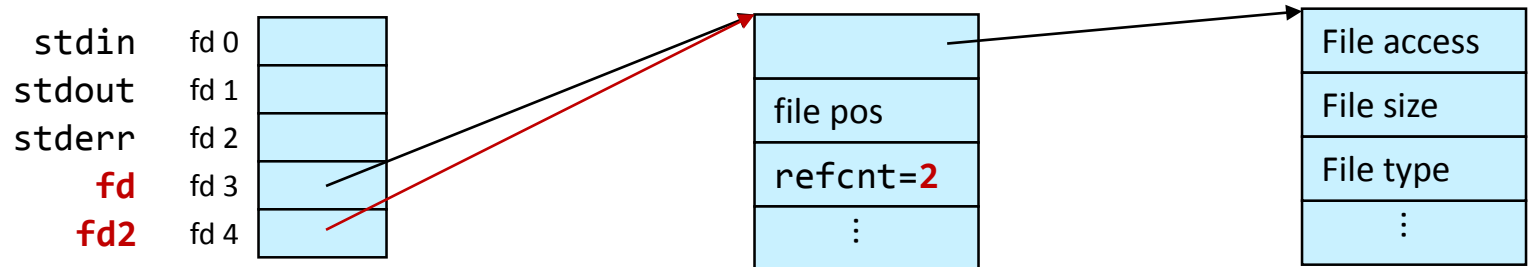
Operation	API	Remarks
Duplicate file descriptor	<code>int dup(int oldfd)</code>	Returned fd points to same entry in open file table as oldfd
Duplicate file descriptor to specific fd	<code>int dup2(int oldfd, int newfd)</code>	Entry in newfd is overwritten with value in oldfd
Retrieve file descriptor of file stream	<code>int fileno(FILE *stream)</code>	

File Descriptor Manipulation: dup()

■ `fd = open("file.txt", O_RDONLY, 0);`




■ `fd2 = dup(fd);`

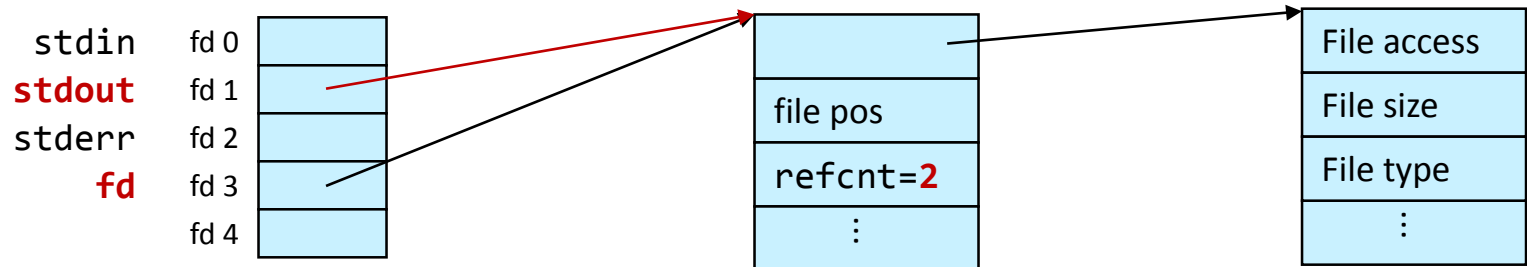


File Descriptor Manipulation: dup2()

■ `fd = open("output.txt", O_WRONLY, 0);`



■ `dup2(fd, STDIO_FILENO);`



Fun with File Descriptors

- What is the output of this program if the input file contains “System Programming”?

```
#include <...>

int main(int argc, char *argv[])
{
    int fd1, fd2, fd3;
    char c1, c2, c3;
    char *fname = argv[1];

    fd1 = open(fname, O_RDONLY, 0);
    fd2 = open(fname, O_RDONLY, 0);
    fd3 = open(fname, O_RDONLY, 0);
    if ((fd1 == -1) || (fd2 == -1) || (fd3 == -1)) {
        fprintf(stderr, "Cannot open input file.\n"); return EXIT_FAILURE;
    }

    dup2(fd2, fd3);

    read(fd1, &c1, 1);
    read(fd2, &c2, 1);
    read(fd3, &c3, 1);

    printf("c1 = %c, c2 = %c, c3 = %c\n", c1, c2, c3);
    return 0;
}
```

fwfd1.c

Fun with File Descriptors

- What are the contents of the generated output file?

```
#include <...>

int main(int argc, char *argv[])
{
    int fd1, fd2, fd3;
    char *fname = argv[1];

    if ((fd1 = open(fname, O_CREAT|O_TRUNC|O_RDWR, S_IRUSR|S_IWUSR)) == -1) {
        perror("Cannot open output file"); return EXIT_FAILURE;
    }

    write(fd1, "CSAP", 4);

    fd3 = open(fname, O_APPEND|O_WRONLY, 0);
    write(fd3, "M1522", 5);

    fd2 = dup(fd1);                // Allocates descriptor
    write(fd2, "SNU", 3);
    write(fd3, "800", 3);

    return 0;
}
```

fwfd2.c

Summary

- -----
- -----
- -----
- -----
- -----
- -----

Module Summary

Most Frequently Used Unix I/O System Calls

```
#include <unistd.h>
```

`$ man -S 2 <func>`

```
#include <fcntl.h>
```

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int      open(const char *pathname, int flags[, mode_t mode]);
```

```
int      creat(const char *pathname, mode_t mode);
```

```
ssize_t read(int fd, void *buf, size_t count);
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

```
off_t    lseek(int fd, off_t offset, int whence);
```

```
int      stat(const char *path, struct stat *buf);
```

```
int      close(int fd);
```

Most Frequently Used Standard I/O System Calls

```
#include <stdio.h>
```

```
$ man -S 3 <func>
```

```
FILE*    fopen(const char *pathname, const char *mode);
```

```
size_t   fread(void *ptr, size_t size, size_t nmemb, FILE *stream);
```

```
size_t   fwrite(const void *ptr, size_t size, size_t nmemb, FILE *stream);
```

```
int       fflush(FILE *stream);
```

```
int       feof(FILE *stream);
```

```
int       ferror(FILE *stream);
```

```
off_t     fseek(FILE *stream, long offset, int whence);
```

```
int       f[get/set]pos(FILE *stream, fpos_t *pos);
```

```
int       fclose(FILE *fp);
```

Most Frequently Used API to Manage Directories

```
#include <sys/types.h>
```

```
$ man -S 3 <func>
```

```
#include <dirent.h>
```

```
DIR*      opendir(const char *name);
```

```
struct dirent *readdir(DIR *dirp);
```

```
int       dirfd(DIR *dirp);
```

```
int       closedir(DIR *dirp);
```


Pros and Cons of Unix I/O

■ Pros

- Unix I/O is the most general and lowest overhead form of I/O.
 - ▶ All other I/O packages are implemented using Unix I/O functions.
- Unix I/O provides functions for accessing file metadata.
- Unix I/O functions are async-signal-safe and can be used safely in signal handlers.

■ Cons

- Dealing with short counts is tricky and error prone.
- Efficient reading of text lines requires some form of buffering, also tricky and error prone.
- Both of these issues are addressed by the standard I/O packages.

Pros and Cons of Standard I/O

■ Pros

- Buffering increases efficiency by decreasing the number of read and write system calls
- Short counts are handled automatically

■ Cons

- Provides no function for accessing file metadata
- Standard I/O functions are not async-signal-safe, and not appropriate for signal handlers.
- Standard I/O is not appropriate for input and output on network sockets
 - ▶ There are poorly documented restrictions on streams that interact badly with restrictions on sockets (CS:APP3e, Sec 10.9)

Choosing I/O Functions

- General rule: use the highest-level I/O functions you can
 - Many C programmers are able to do all of their work using the standard I/O functions
- When to use standard I/O
 - When working with disk or terminal files
- When to use raw Unix I/O
 - Inside signal handlers, because Unix I/O is async-signal-safe.
 - In rare cases when you need absolute highest performance.

Aside: Working with Binary Files

- Binary File Examples
 - Object code, Images (JPEG, GIF)
- Functions you shouldn't use on binary files
 - Line-oriented I/O such as fgets, scanf, printf
 - ▶ Different systems interpret 0x0A ('\n') (newline) differently:
 - Linux and Mac OS X: LF(0x0a) ['\n']
 - HTTP servers & Windows: CR+LF(0x0d 0x0a) ['\r\n']
 - String functions
 - ▶ strlen, strcpy
 - ▶ Interprets byte value 0 (end of string) as special

For Further Information

- The Unix bible:

- W. Richard Stevens & Stephen A. Rago, Advanced Programming in the Unix Environment, 3rd Edition, Addison-Wesley, 2013, ISBN 978-0321637734
 - ▶ Updated from Stevens's 1993 classic text.

